



The effect of extra-sentential context on Speech to Discrete Unit Translation

B235864

7369 words

Master of Science

Speech and Language Processing

School of Philosophy, Psychology and Language
Sciences

University of Edinburgh

2024

Abstract

Contextual information, possibly from previous sentences, is well known to be useful during translation, for resolving ambiguity or translating obscure words. This has been successfully exploited by textual machine translation, and similar success has recently been found in direct speech-to-text translation. More end-to-end speech-to-speech translation is yet to take advantage of these extra-sentential context clues, and it is not clear how well existing methods might apply. This work shows that extra-sentential context can be usefully incorporated into direct speech-to-speech translation, using a method called speech-to-unit translation (S2UT), and a crucial tool for reducing redundancy in the speech input, Adaptive Feature Selection. Results show that S2UT is sensitive to training data choices and parameter choices, but despite that sensitivity, incorporating extra-sentential context yields improvements of up to +5.4 Bleu score over the S2UT baseline.

Acknowledgements

I would like to give thanks to my supervisor Eva Colley, and the rest of the team at CSLT for their guidance and expertise on the topic of speech translation. I would like to offer enormous thanks to the coding help staff at Edinburgh University for helping with some challenging technical issues with FAIRSEQ, and for patiently listening to me explain the details of my project. Finally, I want to thank my family and friends for helping me through this project, and reminding me that there is definitely more to life than Speech to Speech translation.